

Task 2 — Evaluating Chaotic Systems

Ahmed Usama Masoud

Phase-space reconstruction · Lyapunov spectrum · K_2 entropy · Correlation & box-counting dimensions

System under study — Lorenz.

$$\dot{x} = \sigma(y - x), \quad \dot{y} = x(\rho - z) - y, \quad \dot{z} = x \cdot y - \beta \cdot z$$

Classical chaotic parameters $\sigma = 10$, $\rho = 28$, $\beta = 8/3$ were used throughout. The system was integrated with SciPy's `solve_ivp` (RK45) at $dt = 0.01$ over $T = 150$, and the first 20 time units (2000 samples) were discarded as transient before any analysis. The scalar series used for parts A, B (time-series method), C and D is the x-component of the trajectory; part E uses the full 3-D state.

1. Part A — Phase Space Reconstruction

1.1 Concept

When only a scalar time series $x(t)$ is available, Takens' embedding theorem guarantees that a topologically equivalent copy of the attractor can be reconstructed from delay vectors $y(t) = [x(t), x(t+\tau), x(t+2\tau), \dots, x(t+(m-1)\tau)]$. Two parameters must be chosen: the time delay τ (too small $\rightarrow$ over-correlated coordinates; too large $\rightarrow$ decorrelated coordinates) and the embedding dimension m (too small $\rightarrow$ false crossings of the reconstructed attractor).

1.2 Time delay via Average Mutual Information (AMI)

AMI generalises linear autocorrelation to a nonlinear, information-theoretic measure of dependence between x_i and $x_{i+\tau}$. The optimal delay is the first local minimum of $I(\tau)$ as τ increases (Fraser & Swinney, 1986): at this lag the two coordinates share the least redundant information while still being dynamically related.




Figure A.1 — Average Mutual Information vs lag for the Lorenz x-series. Red dashed line marks the first local minimum $\tau =$

Result. The AMI curve drops sharply from lag 0 and reaches its first local minimum at $\tau = 16$ samples (≈ 0.16 time units at $dt = 0.01$). This falls within the expected range for Lorenz (≈ 10 – 17).

1.3 Embedding dimension via False Nearest Neighbours (FNN)

For each candidate dimension d the nearest neighbour of every reconstructed point is tested with the two Kennel criteria; a neighbour that separates sharply when one more coordinate is added is flagged as false. The optimal m is the smallest dimension at which the FNN percentage drops below a chosen threshold (1% here).




Figure A.2 — FNN percentage vs embedding dimension. Green dashed line marks the selected $m = 3$; red dashed line is the 1% threshold.

Result. The FNN percentage collapses from $\approx 100\%$ at $m = 1$ to $\approx 30\%$ at $m = 2$ and essentially to zero at $m = 3$, and stays there for larger dimensions. The optimal embedding is therefore $m = 3$, in perfect agreement with the true state dimension of the Lorenz system.

Settings:

Setting	Value
Series	$x(t)$ from Lorenz, $T = 150$, $dt = 0.01$, transient 20 t.u. discarded
Series length after transient	13 000 samples
AMIconfig	max_lag = 200, n_bins = 32, criterion = first_local_min
FNNConfig	max_dim = 15, R_tol = 10.0, A_tol = 2.0, threshold_percent = 1.0, theiler = 100
Result τ	16 samples
Result m	3
Expected range (Lorenz)	$\tau \sim 10$ – 17 , $m \sim 3$ – 7

2. Part B — Lyapunov Exponent

2.1 Concept

The Lyapunov exponent λ quantifies the average exponential rate of divergence of infinitesimally close trajectories. A system is chaotic when its largest Lyapunov exponent λ_1 is strictly positive (bounded flow with local exponential divergence); $\lambda_1 = 0$ indicates a marginal/periodic case and $\lambda_1 < 0$ indicates convergence to a fixed point. A hyperchaotic system has two or more positive exponents.

2.2 From the ODE — Wolf's Gram–Schmidt Reorthogonalisation (GSR)

With the equations and Jacobian known, `lyapunov_wolf_ode` evolves n orthonormal tangent vectors alongside the trajectory via the variational equation $\dot{V} = J(X(t)) \cdot V$, and performs a QR decomposition every `ortho_steps` integration steps. The diagonal of R gives the local stretching factors along each principal direction; accumulating their logs over time yields the full spectrum ($\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_n$).

Result ($\log_2$ units, bits/time):

λ_1	λ_2	λ_3
+1.2966	-0.0033	-20.9980

Exactly one positive exponent confirms Lorenz is chaotic (not hyperchaotic). λ_2 is close to zero as expected for a continuous flow (the marginal direction along the trajectory itself), and the large negative λ_3 reflects the strong contraction that pulls trajectories onto the 2-D attractor sheet.

Settings — `lyapunov_wolf_ode`:

Setting	Value
ODE / Jacobian	Lorenz ($\sigma = 10$, $\rho = 28$, $\beta = 8/3$)
Initial state x_0	(1, 1, 1)
dt	0.01
transient_steps	1000
n_steps	30 000
Solver	RK45 (scipy solve_ivp)
ortho_steps	20
log_base	2 (bits/time)

2.3 From a time series — Wolf's phase-space tracking (`wolf_mle`)

When only a scalar series is available, the phase space is first reconstructed with the (τ, m) from Part A. From a reference point, `wolf_mle` finds a valid nearest neighbour outside the Theiler window and evolves both forward, tracking their separation $d_k = \|x(t_k) - x'(t_k)\|$ until it exceeds a divergence threshold; a new neighbour is then chosen to keep the separation small while preserving the previous vector's orientation. Averaging $\log_2(d_k/d_{k-1})$ over all tracked segments gives λ_1 in bits per time unit.

Result:

λ_1 (time-series, wolf_mle) = 1.8128 bits/time (200 segments tracked)

λ_1 (ODE, lyapunov_wolf_ode) = 1.2966 bits/time

Comparison. Both estimates are positive and of the same order of magnitude, so the two methods agree qualitatively that the Lorenz system is chaotic. The $\sim 40\%$ discrepancy between the two numbers is the classical time-series-vs-ODE gap: the tracking method is sensitive to the finite series length, to the exact choice of the min/max distance band and evolution cap, and to how faithfully the (τ, m) embedding reproduces the true tangent-space structure. The ODE-based spectrum is the reference value; the widely-cited numerical value in the literature for Lorenz at $p = 28$ is $\lambda_1 \approx 0.906$ nats/time ≈ 1.30 bits/time, which matches the ODE result here to two decimal places.

Settings — wolf_mle:

Setting	Value
Series	same Lorenz x-series as Part A (11 000 samples used)
dt	0.01
τ, m	16, 3 (from Part A)
theiler_window	default = $\tau \cdot m = 48$
min_dist_scale / max_dist_scale	1e-3 / 1e-1 of median pairwise distance
evolve_cap	50
max_replacements	200
angle_weight	0.3
log_base (reported)	2 (bits/time)

3. Part C — Kolmogorov (K_2) Entropy

The Kolmogorov–Sinai entropy measures the average rate of information production (loss of predictability) of the system. By Pesin's identity it is bounded above by the sum of the positive Lyapunov exponents, and it is closely related in concept to the Lyapunov exponent: K_2 measures the average amount of new information generated per unit time by the chaotic dynamics, whereas λ_1 measures the rate at which nearby trajectories diverge. For a low-dimensional chaotic system with a single positive λ (like Lorenz) the two quantities should therefore be of the same order of magnitude.

The computable lower-bound estimate used here comes from the scaling of the correlation sum $C_m(r)$ with the embedding dimension:

$$K_2 \approx (1/\tau) \cdot \ln [C_m(r) / C_{m+1}(r)] \quad (\text{averaged over the scaling region})$$




Figure C.1 — K_2 estimator as a function of $\log(r)$ for the reconstructed Lorenz attractor at $(\tau, m) = (16, 3)$. The red dashed line marks the reported $K_2 = 4.07$ (see comment below).

Comment — this value contains an error.

The K_2 value shown in Figure C.1 ($K_2 \approx 4.07$ nats/time) is not reasonable: since K_2 and the largest Lyapunov exponent measure closely related quantities (K_2 upper-bounded by the sum of the positive Lyapunov exponents, which for Lorenz is just λ_1), K_2 for Lorenz should be close to the Lyapunov exponent value, i.e. near ≈ 1.8 (in the same units), not ≈ 4 . The estimate from a single (τ, m) point at $m = 3$ is unreliable — it picks up curvature at the edges of the scaling region and overestimates the entropy. The correct procedure is to repeat the estimate for a range of embedding dimensions and read the saturation value, as shown next.




Figure C.2 — K_2 estimate as a function of embedding dimension $m \in [2, 8]$.

Result. As m increases the K_2 estimate drops and saturates at $K_2 \approx 1.83$ nats/time — which is the physically meaningful value, and is exactly of the same order as λ_1 from Part B, consistent with the fact that Lorenz has a single positive Lyapunov exponent. The saturation of K_2 with m (rather than growth without bound) is the signature of low-dimensional deterministic chaos, not noise.

Settings:

Setting	Value
Estimator	$K_2 = (1/\tau) \cdot \ln[C_m(r)/C_{m+1}(r)]$ over scaling region
Embedding τ	16 (from Part A)
Embedding dims tested	$m \in [2, 8]$ (single-m panel uses $m = 3$)
Reported single-m K_2 ($m = 3$)	4.07 nats/time ← unreliable, see comment
Saturated K_2 ($m = 7$)	1.83 nats/time ← physically meaningful value

4. Part D — Correlation Dimension (D_2)

D_2 is the Correlation Dimension of an attractor. It is called "correlation" because it is computed from the correlation sum $C(r)$, which measures how many pairs of reconstructed points lie within a distance r of each other:

$$C(r) = (2 / N(N-1)) \cdot \sum_{\{i < j\}} \Theta(r - \|y_i - y_j\|)$$

The Grassberger–Procaccia algorithm then defines D_2 as the small- r scaling exponent, $C(r) \propto r^{D_2}$, and estimates it as the slope of the linear region of $\log C(r)$ vs. $\log r$.




Figure D.1 — Grassberger–Procaccia $\log C(r)$ vs $\log r$ for the Lorenz attractor at $m = 3$. The dashed line is the linear fit; its slope gives $D_2 = 2.056$.




Figure D.2 — D_2 estimate vs embedding dimension $m \in [2, 12]$. The curve saturates quickly, confirming a low-dimensional attractor.

Result. The D_2 curve is flat from $m = 3$ onward and saturates at $D_2 \approx 2.09$, comfortably below the state dimension 3 — as expected for the Lorenz attractor, which is a strange (fractal) attractor with a non-integer dimension slightly above 2. Saturation with m (rather than growing linearly) again confirms low-dimensional deterministic chaos rather than noise.

Settings:

Setting	Value
Algorithm	Grassberger–Procaccia via <code>estimate_correlation_dimension</code>
Embedding τ	16 (from Part A)
Embedding dims tested	$m \in [2, 12]$
Radius range	logarithmically spaced, spanning the attractor's linear-scaling region
Theiler window	default (excludes temporally close pairs)
Reported D_2 ($m = 3$)	2.056
Saturated D_2 ($m = 11$)	2.086

5. Part E — Fractal (Box-Counting) Dimension (D_0)

Part E is the most intuitive of the three dimensions. While Part D looked at pairwise distances between points, Part E asks the much simpler question: "how many small boxes are needed to cover the attractor?" The attractor is covered with a uniform grid of boxes of size ϵ , $N(\epsilon)$ is the number of boxes containing at least one point of the trajectory, and D_0 is defined as the scaling exponent

$$D_0 = \lim_{\epsilon \rightarrow 0} \{ \ln N(\epsilon) / \ln (1/\epsilon) \}$$

It is estimated as the slope of $\ln N(\epsilon)$ vs $\ln(1/\epsilon)$ over the linear-scaling region as ϵ is decreased.




Figure E — Box-counting: $\ln N(\epsilon)$ vs $\ln(1/\epsilon)$ for the full 3-D Lorenz trajectory. The dashed line is the linear fit; its slope gives $D_0 = 1.802$.

Result and comment.

The reported box-counting dimension is $D_0 \approx 1.80$. Numerically this is close to the widely quoted Lorenz value (≈ 2.06 for the true dimension), but it is smaller than the $D_2 \approx 2.09$ obtained in Part D — which contradicts the general ordering $D_0 \geq D_2$. This value therefore contains some error.

Settings:

Setting	Value
Input	Full 3-D Lorenz trajectory (x, y, z), 13 000 samples after transient
Algorithm	estimate_box_dimension
Box sizes ϵ	log-spaced from $\approx$ attractor diameter down to resolution limit
Reported D_0	1.802 (contains finite-N error, see comment)
Reference / D_2 from Part D	≈ 2.09 (D_0 should be $\geq$ this)

6. Conclusion

The five diagnostics tell a consistent story about the Lorenz system. Phase-space reconstruction from a single scalar series recovers exactly the expected embedding ($\tau = 16$, $m = 3$), matching the true state dimension. The full Lyapunov spectrum from the ODE contains exactly one positive exponent ($\lambda_1 \approx 1.30$ bits/time), one near-zero exponent along the flow, and one strongly negative exponent the fingerprint of an ordinary (not hyperchaotic) 3-D dissipative chaotic system. The time-series estimate of λ_1 (≈ 1.81 bits/time) is of the same order and confirms the classification qualitatively. The Kolmogorov entropy saturates at $K_2 \approx 1.83$ nats/time with embedding dimension of the same order as λ_1 , as expected from Pesin's inequality and both the correlation dimension ($D_2 \approx 2.09$) and the box-counting dimension ($D_0 \approx 1.80$, subject to the finite-size caveat noted in Part E) place the attractor in the expected fractal range, slightly above 2 and below the state dimension 3. Taken together one positive Lyapunov exponent, a finite positive entropy, and a non-integer dimension between 2 and 3 these values confirm that Lorenz is a genuine low-dimensional chaotic (not hyperchaotic) attractor.